

# Investigate\_a\_Dataset

September 6, 2022

## 1 Project: Investigate TMDb Movies Dataset

### 1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

## Introduction

#### 1.1.1 Dataset Description

This report was generated from a dataset collection from The Movie Database (TMDb). This is a very thorough dataset of movies release between the years of 1960 and 2015. There are over 10,000 rows of data with characteristics of the film (tagline, genre, etc), financial information, popularity, and when the film was release.

#### 1.1.2 Questions for Analysis

Have movies been getting longer year over year?

What year had the most popular releases?

Does the vote averages reflect similar results as popularity?

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: # Upgrade pandas to use dataframe.explode() function.
#!pip install --upgrade pandas==0.25.0
```

## Data Wrangling Step 1: Gather Data

First, I used pandas to load the dataset.

```
In [3]: df_movies = pd.read_csv('Database_TMDb_movie_data/tmdb-movies.csv')
```

## 1.2 Data Wrangling Step 2: Assess Data

Second, I used several function to get a general understanding of the dataset. I found that there are 10,866 rows and 21 column of data. There is a mixture of numeric and string datatypes and several cells using an object data type had multiple point of thata that were separated by a bar ( | ). On initial glance it looked as if the majority of the omitted data was confined to the object datatype, but after review the statistical inforamtion it appears that missing numerical data simply had a "0" as a placeholder.

```
In [4]: df_movies.shape
```

```
Out[4]: (10866, 21)
```

```
In [5]: df_movies.head(5)
```

```
Out[5]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	<a href="http://www.jurassicworld.com/">http://www.jurassicworld.com/</a>	Colin Trevorrow	
1	<a href="http://www.madmaxmovie.com/">http://www.madmaxmovie.com/</a>	George Miller	
2	<a href="http://www.thedivergentseries.movie/#insurgent">http://www.thedivergentseries.movie/#insurgent</a>	Robert Schwentke	
3	<a href="http://www.starwars.com/films/star-wars-episod...">http://www.starwars.com/films/star-wars-episod...</a>	J.J. Abrams	
4	<a href="http://www.furious7.com/">http://www.furious7.com/</a>	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

		overview	runtime	\
0	Twenty-two years after the events of Jurassic ...		124	
1	An apocalyptic story set in the furthest reach...		120	
2	Beatrice Prior must confront her inner demons ...		119	
3	Thirty years after defeating the Galactic Empi...		136	
4	Deckard Shaw seeks revenge against Dominic Tor...		137	

		genres	\
0	Action Adventure Science Fiction Thriller		
1	Action Adventure Science Fiction Thriller		
2	Adventure Science Fiction Thriller		
3	Action Adventure Science Fiction Fantasy		
4	Action Crime Thriller		

		production_companies	release_date	vote_count	\
0	Universal Studios Amblin Entertainment Legenda...		6/9/15	5562	
1	Village Roadshow Pictures Kennedy Miller Produ...		5/13/15	6185	
2	Summit Entertainment Mandeville Films Red Wago...		3/18/15	2480	
3	Lucasfilm Truenorth Productions Bad Robot		12/15/15	5292	
4	Universal Pictures Original Film Media Rights ...		4/1/15	2947	

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

In [6]: df\_movies.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
```

```

runtime          10866 non-null int64
genres           10843 non-null object
production_companies 9836 non-null object
release_date     10866 non-null object
vote_count       10866 non-null int64
vote_average     10866 non-null float64
release_year     10866 non-null int64
budget_adj       10866 non-null float64
revenue_adj      10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

```
In [7]: df_movies.describe()
```

```

Out[7]:

```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

### 1.3 Data Wrangling Step 3: Cleaning Data

Cleaning the data for this dataset required careful consideration. I started with the easy task of removing the 1 duplicated row, as well as many of the columns that I wouldn't be using. Then explored the main columns that I wanted to explore, release year, runtime, popularity and vote average. I wanted to find where any potential outliers were, as well as if they had a placeholder to hide missing data.

```
In [8]: sum(df_movies.duplicated())
```

```
Out[8]: 1
```

```
In [9]: df_movies.drop_duplicates(inplace=True)
```

```
In [10]: sum(df_movies.duplicated())
```

```
Out[10]: 0
```

```
In [11]: df_movies.drop(['imdb_id', 'budget', 'revenue', 'original_title', 'cast', 'homepage', ''],
                        df_movies.head(10))
```

```
Out[11]:
```

	id	popularity	runtime	release_date	vote_count	vote_average	\
0	135397	32.985763	124	6/9/15	5562	6.5	
1	76341	28.419936	120	5/13/15	6185	7.1	
2	262500	13.112507	119	3/18/15	2480	6.3	
3	140607	11.173104	136	12/15/15	5292	7.5	
4	168259	9.335014	137	4/1/15	2947	7.3	
5	281957	9.110700	156	12/25/15	3929	7.2	
6	87101	8.654359	125	6/23/15	2598	5.8	
7	286217	7.667400	141	9/30/15	4572	7.6	
8	211672	7.404165	91	6/17/15	2893	6.5	
9	150540	6.326804	94	6/9/15	3935	8.0	


	release_year
0	2015
1	2015
2	2015
3	2015
4	2015
5	2015
6	2015
7	2015
8	2015
9	2015

I started with reviewing the release year and found the range to be between 1965 and 2015.

```
In [12]: df_movies['release_year'].unique()
```

```
Out[12]: array([2015, 2014, 1977, 2009, 2010, 1999, 2001, 2008, 2011, 2002, 1994,
                2012, 2003, 1997, 2013, 1985, 2005, 2006, 2004, 1972, 1980, 2007,
                1979, 1984, 1983, 1995, 1992, 1981, 1996, 2000, 1982, 1998, 1989,
                1991, 1988, 1987, 1968, 1974, 1975, 1962, 1964, 1971, 1990, 1961,
                1960, 1976, 1993, 1967, 1963, 1986, 1973, 1970, 1965, 1969, 1978,
                1966])
```

```
In [13]: df_movies['release_year'].hist();
```




This release year histogram tells me that the number of movies made year over year has been increasing year over year with an exponential increase in the mid 2000's

Exploring the popularity column presented me with quite a challenge. It was clear by looking at the distribution and the statistical information that there are big outliers that significantly skew the data to the right. After assessing the quartiles of the data, I came to the conclusion that anything greater than a popularity of 1 is a possible outlier.

With the popularity data now in decimals, I attempted to make the analysis easier to read by multiplying the values by 100, and rounding to the nearest whole number.

I verified my assumptions about my assumptions about any possible outliers by creating another histogram with the updates and then saved it as a new dataframe.

```
In [14]: df_movies['popularity'].hist();
```



```
In [15]: df_movies['popularity'].describe()
```

```
Out[15]: count      10865.000000
         mean         0.646446
         std         1.000231
         min         0.000065
         25%         0.207575
         50%         0.383831
         75%         0.713857
         max         32.985763
         Name: popularity, dtype: float64
```

```
In [16]: pop_possible_outliers = df_movies[(df_movies['popularity'] >= 1) | (df_movies['popularity'] >= 30)]
         pop_possible_outliers['popularity'].count()
```

```
Out[16]: 1854
```

```
In [17]: df_movies['pop_round_100'] = df_movies['popularity'].values
         df_movies['pop_round_100'] = np.trunc(df_movies['pop_round_100'] * 100)
         df_movies.describe()
```

```
Out[17]:
```

	id	popularity	runtime	vote_count	vote_average \
count	10865.000000	10865.000000	10865.000000	10865.000000	10865.000000
mean	66066.374413	0.646446	102.071790	217.399632	5.975012
std	92134.091971	1.000231	31.382701	575.644627	0.935138
min	5.000000	0.000065	0.000000	10.000000	1.500000

25%	10596.000000	0.207575	90.000000	17.000000	5.400000
50%	20662.000000	0.383831	99.000000	38.000000	6.000000
75%	75612.000000	0.713857	111.000000	146.000000	6.600000
max	417859.000000	32.985763	900.000000	9767.000000	9.200000

	release_year	pop_round_100
count	10865.000000	10865.000000
mean	2001.321859	64.146618
std	12.813260	100.022164
min	1960.000000	0.000000
25%	1995.000000	20.000000
50%	2006.000000	38.000000
75%	2011.000000	71.000000
max	2015.000000	3298.000000

```
In [18]: df_pop = df_movies[(df_movies['pop_round_100'] < 100) & (df_movies['pop_round_100'] >=
df_pop.describe())
```

```
Out[18]:
```

	id	popularity	runtime	vote_count	vote_average \
count	9011.000000	9011.000000	9011.000000	9011.000000	9011.000000
mean	67153.272889	0.369880	100.308512	71.060260	5.891544
std	92757.981771	0.235482	30.718330	148.513246	0.939037
min	6.000000	0.010016	0.000000	10.000000	1.500000
25%	11363.500000	0.187407	90.000000	15.000000	5.300000
50%	21948.000000	0.322496	97.000000	29.000000	5.900000
75%	74751.500000	0.513636	109.000000	72.000000	6.500000
max	414419.000000	0.999866	877.000000	4368.000000	8.900000

	release_year	pop_round_100
count	9011.000000	9011.000000
mean	2000.788481	36.489291
std	13.101088	23.549588
min	1960.000000	1.000000
25%	1994.000000	18.000000
50%	2005.000000	32.000000
75%	2011.000000	51.000000
max	2015.000000	99.000000

Cleaning the vote average column included verifying there was no outliers and missing data.

```
In [19]: df_pop['vote_average'].describe()
```

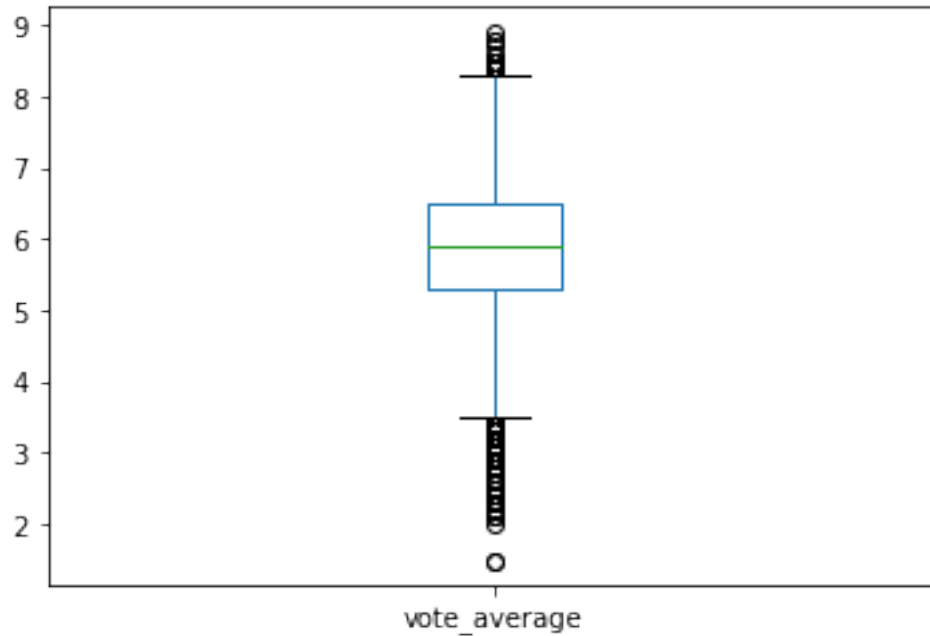
```
Out[19]:
```

count	9011.000000
mean	5.891544
std	0.939037
min	1.500000
25%	5.300000
50%	5.900000




```
75%          6.500000
max          8.900000
Name: vote_average, dtype: float64
```

```
In [20]: df_pop['vote_average'].plot(kind='box');
```




```
In [21]: df_pop['vote_average'].hist();
```



The final cleaning was for the 'runtime' column. I research film lengths I found that the Academy of Motion Picture Arts and Science defines a feature film length is over 40 minutes. Once I removed the lower limit from the dataset, I used three times the standard deviation plus the mean to calculate the upper limit of the data.

This put the film runtime between 40 and 191 minutes.

```
In [22]: df_pop['runtime'].hist();
```



```
In [23]: df_pop['runtime'].describe()
```

```
Out[23]: count      9011.000000
         mean       100.308512
         std        30.718330
         min         0.000000
         25%        90.000000
         50%        97.000000
         75%       109.000000
         max       877.000000
         Name: runtime, dtype: float64
```

```
In [24]: df_feature = df_pop[df_pop['runtime'] > 40]
         df_feature['runtime'].describe()
```

```
Out[24]: count      8779.000000
         mean       102.646429
         std        27.451472
         min        41.000000
         25%        90.000000
         50%        98.000000
         75%       110.000000
         max        877.000000
         Name: runtime, dtype: float64
```

```
In [25]: df_feature = df_feature[df_feature['runtime'] <= 191]
         df_feature.describe()
```

```
Out[25]:
```

	id	popularity	runtime	vote_count	vote_average \
count	8707.000000	8707.000000	8707.000000	8707.000000	8707.000000
mean	65505.695417	0.372821	101.089009	71.991501	5.863225
std	91623.068017	0.235435	18.214575	150.614390	0.928717
min	6.000000	0.010016	41.000000	10.000000	1.500000
25%	11241.500000	0.190222	90.000000	15.000000	5.300000
50%	21183.000000	0.324081	98.000000	29.000000	5.900000
75%	71231.000000	0.516744	109.000000	73.000000	6.500000
max	409696.000000	0.999866	191.000000	4368.000000	8.900000


	release_year	pop_round_100
count	8707.000000	8707.000000
mean	2000.660503	36.784311
std	13.116560	23.544106
min	1960.000000	1.000000
25%	1993.000000	19.000000
50%	2005.000000	32.000000
75%	2011.000000	51.000000
max	2015.000000	99.000000

## Exploratory Data Analysis

### 1.3.1 Research Question 1: Have movies been getting longer?

I wanted to be able to plot the average runtime per year by years. However 55 years was a lot of points on the plot, so I decided to put the 55 years into 5 year ranges. After that I created a function to calculate the average (in this case runtime) per year ranges.

```
In [26]: runtime_hist=df_feature['runtime'].hist()
         runtime_hist.set_ylabel('Movie Counts')
         runtime_hist.set_xlabel('Run Time in Minutes');
```



The distribution of the runtimes is slightly skewed to the right and we can see the vast majority of movies run for about 85 to 100 minutes.


```
In [27]: df_feature['release_year'].min()
```

```
Out[27]: 1960
```

```
In [28]: year_bins = [1960, 1965, 1970, 1975, 1980, 1985, 1990, 1995, 2000, 2005, 2010, 2016]
         year_range = ['1960-64', '1965-69', '1970-74', '1975-79', '1980-84', '1985-89', '1990-94', '1995-99', '2000-04', '2005-09', '2010-16']
```

```
         df_feature['yearRange'] = pd.cut(df_feature['release_year'], bins=year_bins, labels=year_range)
```

```
In [29]: df_feature.groupby('yearRange')['id'].count().plot(kind='bar').set_ylabel('Count');
```




This chart tells me that the amount of movies made since the last 1960's has been increasing and there was an exponential increase between the early and late 2000's.

```
In [30]: """This function calculates the average variable per year range."""
def avg_per_year(x):
    return df_feature.groupby(df_feature.yearRange)[x].mean()

In [31]: plt.figure(figsize=(10,5))
plt.xlabel('Years', fontsize=13)
plt.ylabel('Average Runtime', fontsize=13)
plt.title('Average Runtime Per Year', fontsize=15)

plt.plot(year_range, avg_per_year(['runtime']));
```




The lineplot shows that despite my initial assumption that movies have been getting longer year over year, it is actually the opposite. The average runtime per movie in early 1960's was about 113 minutes, and by the early 2010's the runtime was 17 minutes less at about 96 minutes.

### 1.3.2 Research Question 2 : What year had the most popular releases?


For my next question, I wanted to explore whether or not one year in particular had more popular movies than the others. I created boxplots so I could view all of the statistical information about the popularity by the year range. Next, I created a bar chart to show the average popularity per year range with a line referencing the overall average popularity.

```
In [32]: df_feature.boxplot(column=['pop_round_100'],by=['yearRange'],figsize=(15,10), showmeans=True)
plt.title('Popularity Statistics by Year Range', fontsize =15)
plt.xlabel('Years', fontsize=13);
```



Viewing boxplots give a lot of helpful insights about the spread of popularity by year. We can see that the median throughout the years stays within about 10 points of each other. We can also clearly see that the early 1990's has the most popular releases not only with the highest average, but it has the highest overall inner quartile range.

```
In [33]: plt.figure(figsize=(10,5))
plt.xlabel('Years', fontsize=13)
plt.ylabel('Popularity', fontsize=13)
plt.title('Average Popularity Per Year', fontsize =15)
plt.axhline(y=(df_feature['pop_round_100'].mean()),linewidth=2, color='k', linestyle='--')
plt.bar(year_range, df_feature.groupby(df_feature.yearRange)['pop_round_100'].mean().values)
```




While the barchart shows very similar information as the boxplots it is easier to read the most popular year range as 1990-1994. It is also interesting to see the overall popularities compared to the year ranges.

### 1.3.3 Research Question 3 : Does the vote averages reflect similar results as popularity?

In my final question I wanted to explore if the voting averages would be a good indicator of overall popularity.


```
In [34]: df_feature['vote_average'].hist();
```





I stared by looking at the distribution, which told me that the vast majority of movies had an average vote between 5 and 7.

```
In [35]: avg_per_year('vote_average').plot(kind='barh', figsize=(5,5));
```



Next I looked at the data broken out by years. This was very interesting as it shows that 1970-1974 had the highest vote average, and the other years didn't have a lot of variation in their averages.

```
In [36]: plt.figure(figsize=(10,5))
```


```
x_axis = np.arange(len(year_range))
bar1_width = x_axis - 0.2
bar2_width = x_axis + 0.2
```

```
y1_pop= avg_per_year('pop_round_100')
y2_votes=avg_per_year('vote_average')
```

```
plt.bar(bar1_width, y1_pop, 0.4, label='Popularity')
plt.bar(bar2_width, y2_votes, 0.4, label='Vote Average')
```

```
plt.xlabel('Years', fontsize=15)
plt.xticks(x_axis, year_range)
```

```
plt.legend();
```



Finally, I plotted the two points together. This data shows that the vote average is not an indicator of the overall popularity.

## ## Conclusions

This data has lead to me to very interesting conclusions. As I was looking into the average runtimes per year my expectation was the opposite of what the data outlined. I assumed that movies overall were getting longer, but according to the data movies between 1960 and 2015 are atually getting shorter on average.

However, there are many considerations that prevent me from making this into a definitive statement. The data only has movies up until 2015, therefore we cannot make draw any concusion after that point. There are also outliers to consider. In my own data cleanup I used three times the standard devation plus the mean to remove outliers, but with a more thorough understanding of the dataset we could find that it was not the best way to handle those outliers.

The popularity graphs show that the 1990's had the most popular releases, and also that there isn't a huge variation on movie popularity in this dataset. With the overall average at 36.8, most of the years were within 5 points of the overall average. We also learned that the average votes does is not an indicator of overall popularity.

I once again cannot make this into a definitive statement because of limitation in the dataset as well as the assumptions I made in handling the outliers. I couldn't find any information about how popularity was measured which is why I tried to a second data

point (vote average) to potentially have a less subjective measurement. It also made it difficult when working on the outliers. I assumed that since 83% of the popularity measurement was less than 1 and greater than 0.01 anything outside of this range could be considered an outlier.

```
In [37]: from subprocess import call  
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[37]: 0
```